

docker build -t spark .

[+] Building 57.0s (10/10) FINISHED

=> [internal] load build definition from Dockerfile

0.4s

=> => transferring dockerfile: 32B

0.0s

=> [internal] load .dockerignore

0.6s

=> => transferring context: 2B

0.0s

=> [internal] load metadata for docker.io/datamechanics/spark:3.1-latest

1.6s

=> [1/5] FROM docker.io/datamechanics/spark:3.1-

latest@sha256:bc32f9839a2b1030d1f20b79725b2058258032d2cfc88e287ed4ca8c91b75124

0.0s

=> [internal] load build context

0.2s

=> => transferring context: 630B

0.0s

=> CACHED [2/5] WORKDIR /opt/application

0.0s

=> CACHED [3/5] RUN pip3 install pyspark

0.0s

=> [4/5] COPY main.py .

0.6s

=> [5/5] COPY data.txt .

43.5s

=> exporting to image

10.3s

=> => exporting layers

9.0s

=> => writing image sha256:0621b8b2875cc49106b87a5efbbe90ff8cf47a0421105884c262b64f17efea47

0.0s

=> => naming to docker.io/library/spark

docker run spark driver local:///opt/application/main.py

Unsetting extraneous env vars (UTC): 11:51:39

Finished unsetting extraneous env vars (UTC): 11:51:39

```
++ id -u
+ myuid=185
++ id -g
+ mygid=0
+ set +e
++ getent passwd 185
+ uidentry=
+ set -e
+ '[' -z '' ']'
+ '[' -w /etc/passwd ']'
+ echo '185:x:185:0:anonymous uid:/opt/spark:/bin/false'
+ SPARK_CLASSPATH=':/opt/spark/jars/*'
+ env
+ grep SPARK_JAVA_OPT_
+ sort -t_ -k4 -n
+ sed 's/[^=]*=\\(.*)\\1/g'
+ readarray -t SPARK_EXECUTOR_JAVA_OPTS
+ '[' -n '' ']'
+ '[' -z ']'
+ '[' -z ']'
+ '[' -n '' ']'
+ '[' -z ']'
+ '[' -z ']'
+ '[' -z x ']'
+ SPARK_CLASSPATH='/opt/spark/conf:/opt/spark/jars/*'
+ case "$1" in
```

+ shift 1

+ CMD=("\$SPARK_HOME/bin/spark-submit" --conf
"spark.driver.bindAddress=\$SPARK_DRIVER_BIND_ADDRESS" --deploy-mode client "\$@")

+ exec /usr/bin/tini -s -- /opt/spark/bin/spark-submit --conf spark.driver.bindAddress= --deploy-mode client
local:///opt/application/main.py

22/12/19 11:51:41 WARN NativeCodeLoader: Unable to load native-hadoop library for your platform...
using builtin-java classes where applicable

Using Spark's default log4j profile: org/apache/spark/log4j-defaults.properties

22/12/19 11:51:42 INFO SparkContext: Running Spark version 3.1.3

22/12/19 11:51:42 INFO ResourceUtils:

=====

22/12/19 11:51:42 INFO ResourceUtils: No custom resources configured for spark.driver.

22/12/19 11:51:42 INFO ResourceUtils:

=====

22/12/19 11:51:42 INFO SparkContext: Submitted application: Python Spark SQL basic example

22/12/19 11:51:42 INFO ResourceProfile: Default ResourceProfile created, executor resources: Map(cores -
> name: cores, amount: 1, script: , vendor: , memory -> name: memory, amount: 1024, script: ,
vendor: , offHeap -> name: offHeap, amount: 0, script: , vendor:), task resources: Map(cpus -> name: cpus,
amount: 1.0)

22/12/19 11:51:42 INFO ResourceProfile: Limiting resource is cpu

22/12/19 11:51:42 INFO ResourceProfileManager: Added ResourceProfile id: 0

22/12/19 11:51:42 INFO SecurityManager: Changing view acls to: 185

22/12/19 11:51:42 INFO SecurityManager: Changing modify acls to: 185

22/12/19 11:51:42 INFO SecurityManager: Changing view acls groups to:

22/12/19 11:51:42 INFO SecurityManager: Changing modify acls groups to:

22/12/19 11:51:42 INFO SecurityManager: SecurityManager: authentication disabled; ui acls disabled; users
with view permissions: Set(185); groups with view permissions: Set(); users with modify p
ermissions: Set(185); groups with modify permissions: Set()

22/12/19 11:51:42 INFO Utils: Successfully started service 'sparkDriver' on port 46155.

22/12/19 11:51:42 INFO SparkEnv: Registering MapOutputTracker

22/12/19 11:51:42 INFO SparkEnv: Registering BlockManagerMaster

22/12/19 11:51:42 INFO BlockManagerMasterEndpoint: Using
org.apache.spark.storage.DefaultTopologyMapper for getting topology information

22/12/19 11:51:42 INFO BlockManagerMasterEndpoint: BlockManagerMasterEndpoint up

22/12/19 11:51:42 INFO SparkEnv: Registering BlockManagerMasterHeartbeat

22/12/19 11:51:42 INFO DiskBlockManager: Created local directory at /tmp/blockmgr-7c27c3dd-b3c0-40fd-97c3-f3186e8ae232

22/12/19 11:51:42 INFO MemoryStore: MemoryStore started with capacity 366.3 MiB

22/12/19 11:51:42 INFO SparkEnv: Registering OutputCommitCoordinator

22/12/19 11:51:43 INFO Utils: Successfully started service 'SparkUI' on port 4040.

22/12/19 11:51:43 INFO SparkUI: Bound SparkUI to 0.0.0.0, and started at http://d0a278a3161b:4040

22/12/19 11:51:43 INFO Executor: Starting executor ID driver on host d0a278a3161b

22/12/19 11:51:43 INFO Utils: Successfully started service 'org.apache.spark.network.netty.NettyBlockTransferService' on port 43681.

22/12/19 11:51:43 INFO NettyBlockTransferService: Server created on d0a278a3161b:43681

22/12/19 11:51:43 INFO BlockManager: Using org.apache.spark.storage.RandomBlockReplicationPolicy for block replication policy

22/12/19 11:51:43 INFO BlockManagerMaster: Registering BlockManager BlockManagerId(driver, d0a278a3161b, 43681, None)

22/12/19 11:51:43 INFO BlockManagerMasterEndpoint: Registering block manager d0a278a3161b:43681 with 366.3 MiB RAM, BlockManagerId(driver, d0a278a3161b, 43681, None)

22/12/19 11:51:43 INFO BlockManagerMaster: Registered BlockManager BlockManagerId(driver, d0a278a3161b, 43681, None)

22/12/19 11:51:43 INFO BlockManager: Initialized BlockManager: BlockManagerId(driver, d0a278a3161b, 43681, None)

22/12/19 11:51:44 INFO SharedState: Setting hive.metastore.warehouse.dir ('null') to the value of spark.sql.warehouse.dir ('file:/opt/application/spark-warehouse').

22/12/19 11:51:44 INFO SharedState: Warehouse path is 'file:/opt/application/spark-warehouse'.

22/12/19 11:51:45 INFO MemoryStore: Block broadcast_0 stored as values in memory (estimated size 294.1 KiB, free 366.0 MiB)

22/12/19 11:51:45 INFO MemoryStore: Block broadcast_0_piece0 stored as bytes in memory (estimated size 27.0 KiB, free 366.0 MiB)

22/12/19 11:51:45 INFO BlockManagerInfo: Added broadcast_0_piece0 in memory on d0a278a3161b:43681 (size: 27.0 KiB, free: 366.3 MiB)

22/12/19 11:51:45 INFO SparkContext: Created broadcast 0 from textFile at NativeMethodAccessorImpl.java:0

22/12/19 11:51:45 INFO FileInputFormat: Total input files to process : 1

22/12/19 11:51:45 INFO SparkContext: Starting job: takeOrdered at /opt/application/main.py:15

22/12/19 11:51:45 INFO DAGScheduler: Got job 0 (takeOrdered at /opt/application/main.py:15) with 100 output partitions

22/12/19 11:51:45 INFO DAGScheduler: Final stage: ResultStage 0 (takeOrdered at /opt/application/main.py:15)

22/12/19 11:51:45 INFO DAGScheduler: Parents of final stage: List()

22/12/19 11:51:45 INFO DAGScheduler: Missing parents: List()

22/12/19 11:51:45 INFO DAGScheduler: Submitting ResultStage 0 (PythonRDD[2] at takeOrdered at /opt/application/main.py:15), which has no missing parents

22/12/19 11:51:46 INFO MemoryStore: Block broadcast_1 stored as values in memory (estimated size 8.5 KiB, free 366.0 MiB)

22/12/19 11:51:46 INFO MemoryStore: Block broadcast_1_piece0 stored as bytes in memory (estimated size 5.0 KiB, free 366.0 MiB)

22/12/19 11:51:46 INFO BlockManagerInfo: Added broadcast_1_piece0 in memory on d0a278a3161b:43681 (size: 5.0 KiB, free: 366.3 MiB)

22/12/19 11:51:46 INFO SparkContext: Created broadcast 1 from broadcast at DAGScheduler.scala:1433

22/12/19 11:51:46 INFO DAGScheduler: Submitting 100 missing tasks from ResultStage 0 (PythonRDD[2] at takeOrdered at /opt/application/main.py:15) (first 15 tasks are for partitions Vector(0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14))

22/12/19 11:51:46 INFO TaskSchedulerImpl: Adding task set 0.0 with 100 tasks resource profile 0

22/12/19 11:51:46 INFO TaskSetManager: Starting task 0.0 in stage 0.0 (TID 0) (d0a278a3161b, executor driver, partition 0, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:46 INFO TaskSetManager: Starting task 1.0 in stage 0.0 (TID 1) (d0a278a3161b, executor driver, partition 1, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:46 INFO TaskSetManager: Starting task 2.0 in stage 0.0 (TID 2) (d0a278a3161b, executor driver, partition 2, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:46 INFO TaskSetManager: Starting task 3.0 in stage 0.0 (TID 3) (d0a278a3161b, executor driver, partition 3, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:46 INFO TaskSetManager: Starting task 4.0 in stage 0.0 (TID 4) (d0a278a3161b, executor driver, partition 4, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:46 INFO TaskSetManager: Starting task 5.0 in stage 0.0 (TID 5) (d0a278a3161b, executor driver, partition 5, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:46 INFO TaskSetManager: Starting task 6.0 in stage 0.0 (TID 6) (d0a278a3161b, executor driver, partition 6, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:46 INFO TaskSetManager: Starting task 7.0 in stage 0.0 (TID 7) (d0a278a3161b, executor driver, partition 7, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:46 INFO TaskSetManager: Starting task 8.0 in stage 0.0 (TID 8) (d0a278a3161b, executor driver, partition 8, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:46 INFO TaskSetManager: Starting task 9.0 in stage 0.0 (TID 9) (d0a278a3161b, executor driver, partition 9, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:46 INFO TaskSetManager: Starting task 10.0 in stage 0.0 (TID 10) (d0a278a3161b, executor driver, partition 10, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:46 INFO TaskSetManager: Starting task 11.0 in stage 0.0 (TID 11) (d0a278a3161b, executor driver, partition 11, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:46 INFO Executor: Running task 5.0 in stage 0.0 (TID 5)

22/12/19 11:51:46 INFO Executor: Running task 8.0 in stage 0.0 (TID 8)

22/12/19 11:51:46 INFO Executor: Running task 0.0 in stage 0.0 (TID 0)

22/12/19 11:51:46 INFO Executor: Running task 10.0 in stage 0.0 (TID 10)

22/12/19 11:51:46 INFO Executor: Running task 1.0 in stage 0.0 (TID 1)

22/12/19 11:51:46 INFO Executor: Running task 9.0 in stage 0.0 (TID 9)

22/12/19 11:51:46 INFO Executor: Running task 6.0 in stage 0.0 (TID 6)

22/12/19 11:51:46 INFO Executor: Running task 11.0 in stage 0.0 (TID 11)

22/12/19 11:51:46 INFO Executor: Running task 4.0 in stage 0.0 (TID 4)

22/12/19 11:51:46 INFO Executor: Running task 3.0 in stage 0.0 (TID 3)

22/12/19 11:51:46 INFO Executor: Running task 2.0 in stage 0.0 (TID 2)

22/12/19 11:51:46 INFO Executor: Running task 7.0 in stage 0.0 (TID 7)

22/12/19 11:51:47 INFO HadoopRDD: Input split: file:/opt/application/data.txt:174643560+19404840

22/12/19 11:51:47 INFO HadoopRDD: Input split: file:/opt/application/data.txt:38809680+19404840

22/12/19 11:51:47 INFO HadoopRDD: Input split: file:/opt/application/data.txt:0+19404840

22/12/19 11:51:47 INFO HadoopRDD: Input split: file:/opt/application/data.txt:194048400+19404840

22/12/19 11:51:47 INFO HadoopRDD: Input split: file:/opt/application/data.txt:155238720+19404840

22/12/19 11:51:47 INFO HadoopRDD: Input split: file:/opt/application/data.txt:77619360+19404840

22/12/19 11:51:47 INFO HadoopRDD: Input split: file:/opt/application/data.txt:19404840+19404840

22/12/19 11:51:47 INFO HadoopRDD: Input split: file:/opt/application/data.txt:213453240+19404840

22/12/19 11:51:47 INFO HadoopRDD: Input split: file:/opt/application/data.txt:135833880+19404840

22/12/19 11:51:47 INFO HadoopRDD: Input split: file:/opt/application/data.txt:58214520+19404840

22/12/19 11:51:47 INFO HadoopRDD: Input split: file:/opt/application/data.txt:116429040+19404840

22/12/19 11:51:47 INFO HadoopRDD: Input split: file:/opt/application/data.txt:97024200+19404840

22/12/19 11:51:53 INFO PythonRunner: Times: total = 6607, boot = 415, init = 102, finish = 6090

22/12/19 11:51:53 INFO PythonRunner: Times: total = 6654, boot = 384, init = 65, finish = 6205

22/12/19 11:51:53 INFO Executor: Finished task 4.0 in stage 0.0 (TID 4). 2202 bytes result sent to driver

22/12/19 11:51:53 INFO Executor: Finished task 11.0 in stage 0.0 (TID 11). 2204 bytes result sent to driver

22/12/19 11:51:53 INFO TaskSetManager: Starting task 12.0 in stage 0.0 (TID 12) (d0a278a3161b, executor driver, partition 12, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:53 INFO Executor: Running task 12.0 in stage 0.0 (TID 12)

22/12/19 11:51:53 INFO TaskSetManager: Starting task 13.0 in stage 0.0 (TID 13) (d0a278a3161b, executor driver, partition 13, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:53 INFO HadoopRDD: Input split: file:/opt/application/data.txt:232858080+19404840

22/12/19 11:51:54 INFO PythonRunner: Times: total = 6760, boot = 412, init = 76, finish = 6272

22/12/19 11:51:54 INFO Executor: Running task 13.0 in stage 0.0 (TID 13)

22/12/19 11:51:54 INFO HadoopRDD: Input split: file:/opt/application/data.txt:252262920+19404840

22/12/19 11:51:54 INFO TaskSetManager: Finished task 4.0 in stage 0.0 (TID 4) in 7829 ms on d0a278a3161b (executor driver) (1/100)

22/12/19 11:51:54 INFO TaskSetManager: Finished task 11.0 in stage 0.0 (TID 11) in 7855 ms on d0a278a3161b (executor driver) (2/100)

22/12/19 11:51:54 INFO PythonRunner: Times: total = 6831, boot = 401, init = 57, finish = 6373

22/12/19 11:51:54 INFO PythonRunner: Times: total = 6851, boot = 390, init = 61, finish = 6400

22/12/19 11:51:54 INFO Executor: Finished task 10.0 in stage 0.0 (TID 10). 2205 bytes result sent to driver

22/12/19 11:51:54 INFO Executor: Finished task 3.0 in stage 0.0 (TID 3). 2208 bytes result sent to driver

22/12/19 11:51:54 INFO TaskSetManager: Starting task 14.0 in stage 0.0 (TID 14) (d0a278a3161b, executor driver, partition 14, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:54 INFO Executor: Running task 14.0 in stage 0.0 (TID 14)

22/12/19 11:51:54 INFO HadoopRDD: Input split: file:/opt/application/data.txt:271667760+19404840

22/12/19 11:51:54 INFO TaskSetManager: Starting task 15.0 in stage 0.0 (TID 15) (d0a278a3161b, executor driver, partition 15, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:54 INFO Executor: Running task 15.0 in stage 0.0 (TID 15)

22/12/19 11:51:54 INFO TaskSetManager: Finished task 3.0 in stage 0.0 (TID 3) in 8003 ms on d0a278a3161b (executor driver) (3/100)

22/12/19 11:51:54 INFO TaskSetManager: Finished task 10.0 in stage 0.0 (TID 10) in 7991 ms on d0a278a3161b (executor driver) (4/100)

22/12/19 11:51:54 INFO HadoopRDD: Input split: file:/opt/application/data.txt:291072600+19404840

22/12/19 11:51:54 INFO PythonAccumulatorV2: Connected to AccumulatorServer at host: 127.0.0.1 port: 60637

22/12/19 11:51:54 INFO Executor: Finished task 9.0 in stage 0.0 (TID 9). 2204 bytes result sent to driver

22/12/19 11:51:54 INFO TaskSetManager: Starting task 16.0 in stage 0.0 (TID 16) (d0a278a3161b, executor driver, partition 16, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:54 INFO TaskSetManager: Finished task 9.0 in stage 0.0 (TID 9) in 8110 ms on d0a278a3161b (executor driver) (5/100)

22/12/19 11:51:54 INFO Executor: Running task 16.0 in stage 0.0 (TID 16)

22/12/19 11:51:54 INFO HadoopRDD: Input split: file:/opt/application/data.txt:310477440+19404840

22/12/19 11:51:54 INFO PythonRunner: Times: total = 7127, boot = 424, init = 60, finish = 6643

22/12/19 11:51:54 INFO Executor: Finished task 5.0 in stage 0.0 (TID 5). 2188 bytes result sent to driver

22/12/19 11:51:54 INFO TaskSetManager: Starting task 17.0 in stage 0.0 (TID 17) (d0a278a3161b, executor driver, partition 17, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:54 INFO Executor: Running task 17.0 in stage 0.0 (TID 17)

22/12/19 11:51:54 INFO TaskSetManager: Finished task 5.0 in stage 0.0 (TID 5) in 8243 ms on d0a278a3161b (executor driver) (6/100)

22/12/19 11:51:54 INFO HadoopRDD: Input split: file:/opt/application/data.txt:329882280+19404840

22/12/19 11:51:54 INFO PythonRunner: Times: total = 7198, boot = 393, init = 56, finish = 6749

22/12/19 11:51:54 INFO Executor: Finished task 8.0 in stage 0.0 (TID 8). 2191 bytes result sent to driver

22/12/19 11:51:54 INFO TaskSetManager: Starting task 18.0 in stage 0.0 (TID 18) (d0a278a3161b, executor driver, partition 18, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:54 INFO TaskSetManager: Finished task 8.0 in stage 0.0 (TID 8) in 8303 ms on d0a278a3161b (executor driver) (7/100)

22/12/19 11:51:54 INFO Executor: Running task 18.0 in stage 0.0 (TID 18)

22/12/19 11:51:54 INFO HadoopRDD: Input split: file:/opt/application/data.txt:349287120+19404840

22/12/19 11:51:54 INFO PythonRunner: Times: total = 7508, boot = 397, init = 58, finish = 7053

22/12/19 11:51:54 INFO Executor: Finished task 2.0 in stage 0.0 (TID 2). 2215 bytes result sent to driver

22/12/19 11:51:54 INFO TaskSetManager: Starting task 19.0 in stage 0.0 (TID 19) (d0a278a3161b, executor driver, partition 19, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:54 INFO Executor: Running task 19.0 in stage 0.0 (TID 19)

22/12/19 11:51:54 INFO TaskSetManager: Finished task 2.0 in stage 0.0 (TID 2) in 8678 ms on d0a278a3161b (executor driver) (8/100)

22/12/19 11:51:54 INFO HadoopRDD: Input split: file:/opt/application/data.txt:368691960+19404840

22/12/19 11:51:58 INFO PythonRunner: Times: total = 11358, boot = 449, init = 73, finish = 10836

22/12/19 11:51:58 INFO Executor: Finished task 1.0 in stage 0.0 (TID 1). 2193 bytes result sent to driver

22/12/19 11:51:58 INFO TaskSetManager: Starting task 20.0 in stage 0.0 (TID 20) (d0a278a3161b, executor driver, partition 20, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:58 INFO Executor: Running task 20.0 in stage 0.0 (TID 20)

22/12/19 11:51:58 INFO TaskSetManager: Finished task 1.0 in stage 0.0 (TID 1) in 12452 ms on d0a278a3161b (executor driver) (9/100)

22/12/19 11:51:58 INFO HadoopRDD: Input split: file:/opt/application/data.txt:388096800+19404840

22/12/19 11:51:58 INFO PythonRunner: Times: total = 11421, boot = 406, init = 45, finish = 10970

22/12/19 11:51:58 INFO Executor: Finished task 7.0 in stage 0.0 (TID 7). 2198 bytes result sent to driver

22/12/19 11:51:58 INFO TaskSetManager: Starting task 21.0 in stage 0.0 (TID 21) (d0a278a3161b, executor driver, partition 21, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:51:58 INFO TaskSetManager: Finished task 7.0 in stage 0.0 (TID 7) in 12547 ms on d0a278a3161b (executor driver) (10/100)

22/12/19 11:51:58 INFO Executor: Running task 21.0 in stage 0.0 (TID 21)

22/12/19 11:51:58 INFO HadoopRDD: Input split: file:/opt/application/data.txt:407501640+19404840

22/12/19 11:52:00 INFO PythonRunner: Times: total = 12854, boot = 470, init = 58, finish = 12326

22/12/19 11:52:00 INFO Executor: Finished task 6.0 in stage 0.0 (TID 6). 2211 bytes result sent to driver

22/12/19 11:52:00 INFO TaskSetManager: Starting task 22.0 in stage 0.0 (TID 22) (d0a278a3161b, executor driver, partition 22, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:00 INFO Executor: Running task 22.0 in stage 0.0 (TID 22)

22/12/19 11:52:00 INFO TaskSetManager: Finished task 6.0 in stage 0.0 (TID 6) in 13942 ms on d0a278a3161b (executor driver) (11/100)

22/12/19 11:52:00 INFO HadoopRDD: Input split: file:/opt/application/data.txt:426906480+19404840

22/12/19 11:52:00 INFO PythonRunner: Times: total = 13112, boot = 457, init = 69, finish = 12586

22/12/19 11:52:00 INFO Executor: Finished task 0.0 in stage 0.0 (TID 0). 2194 bytes result sent to driver

22/12/19 11:52:00 INFO TaskSetManager: Starting task 23.0 in stage 0.0 (TID 23) (d0a278a3161b, executor driver, partition 23, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:00 INFO TaskSetManager: Finished task 0.0 in stage 0.0 (TID 0) in 14237 ms on d0a278a3161b (executor driver) (12/100)

22/12/19 11:52:00 INFO Executor: Running task 23.0 in stage 0.0 (TID 23)

22/12/19 11:52:00 INFO HadoopRDD: Input split: file:/opt/application/data.txt:446311320+19404840

22/12/19 11:52:00 INFO PythonRunner: Times: total = 6089, boot = -27, init = 58, finish = 6058

22/12/19 11:52:00 INFO Executor: Finished task 17.0 in stage 0.0 (TID 17). 2204 bytes result sent to driver

22/12/19 11:52:00 INFO TaskSetManager: Starting task 24.0 in stage 0.0 (TID 24) (d0a278a3161b, executor driver, partition 24, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:00 INFO Executor: Running task 24.0 in stage 0.0 (TID 24)

22/12/19 11:52:00 INFO TaskSetManager: Finished task 17.0 in stage 0.0 (TID 17) in 6111 ms on d0a278a3161b (executor driver) (13/100)

22/12/19 11:52:00 INFO HadoopRDD: Input split: file:/opt/application/data.txt:465716160+19404840

22/12/19 11:52:00 INFO PythonRunner: Times: total = 6106, boot = -151, init = 161, finish = 6096

22/12/19 11:52:00 INFO Executor: Finished task 18.0 in stage 0.0 (TID 18). 2198 bytes result sent to driver

22/12/19 11:52:00 INFO TaskSetManager: Starting task 25.0 in stage 0.0 (TID 25) (d0a278a3161b, executor driver, partition 25, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:00 INFO Executor: Running task 25.0 in stage 0.0 (TID 25)

22/12/19 11:52:00 INFO TaskSetManager: Finished task 18.0 in stage 0.0 (TID 18) in 6407 ms on d0a278a3161b (executor driver) (14/100)

22/12/19 11:52:00 INFO HadoopRDD: Input split: file:/opt/application/data.txt:485121000+19404840

22/12/19 11:52:01 INFO PythonRunner: Times: total = 7158, boot = 1176, init = 0, finish = 5982

22/12/19 11:52:01 INFO Executor: Finished task 14.0 in stage 0.0 (TID 14). 2199 bytes result sent to driver

22/12/19 11:52:01 INFO TaskSetManager: Starting task 26.0 in stage 0.0 (TID 26) (d0a278a3161b, executor driver, partition 26, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:01 INFO Executor: Running task 26.0 in stage 0.0 (TID 26)

22/12/19 11:52:01 INFO TaskSetManager: Finished task 14.0 in stage 0.0 (TID 14) in 7178 ms on d0a278a3161b (executor driver) (15/100)

22/12/19 11:52:01 INFO HadoopRDD: Input split: file:/opt/application/data.txt:504525840+19404840

22/12/19 11:52:02 INFO PythonRunner: Times: total = 7752, boot = -169, init = 192, finish = 7729

22/12/19 11:52:02 INFO Executor: Finished task 16.0 in stage 0.0 (TID 16). 2204 bytes result sent to driver

22/12/19 11:52:02 INFO TaskSetManager: Starting task 27.0 in stage 0.0 (TID 27) (d0a278a3161b, executor driver, partition 27, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:02 INFO TaskSetManager: Finished task 16.0 in stage 0.0 (TID 16) in 7771 ms on d0a278a3161b (executor driver) (16/100)

22/12/19 11:52:02 INFO Executor: Running task 27.0 in stage 0.0 (TID 27)

22/12/19 11:52:02 INFO HadoopRDD: Input split: file:/opt/application/data.txt:523930680+19404840

22/12/19 11:52:02 INFO PythonRunner: Times: total = 8124, boot = 1281, init = 0, finish = 6843

22/12/19 11:52:02 INFO Executor: Finished task 13.0 in stage 0.0 (TID 13). 2236 bytes result sent to driver

22/12/19 11:52:02 INFO TaskSetManager: Starting task 28.0 in stage 0.0 (TID 28) (d0a278a3161b, executor driver, partition 28, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:02 INFO TaskSetManager: Finished task 13.0 in stage 0.0 (TID 13) in 8205 ms on d0a278a3161b (executor driver) (17/100)

22/12/19 11:52:02 INFO Executor: Running task 28.0 in stage 0.0 (TID 28)

22/12/19 11:52:02 INFO HadoopRDD: Input split: file:/opt/application/data.txt:543335520+19404840

22/12/19 11:52:02 INFO PythonRunner: Times: total = 8322, boot = 1350, init = 1, finish = 6971

22/12/19 11:52:02 INFO Executor: Finished task 12.0 in stage 0.0 (TID 12). 2201 bytes result sent to driver

22/12/19 11:52:02 INFO TaskSetManager: Starting task 29.0 in stage 0.0 (TID 29) (d0a278a3161b, executor driver, partition 29, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:02 INFO TaskSetManager: Finished task 12.0 in stage 0.0 (TID 12) in 8402 ms on d0a278a3161b (executor driver) (18/100)

22/12/19 11:52:02 INFO Executor: Running task 29.0 in stage 0.0 (TID 29)

22/12/19 11:52:02 INFO HadoopRDD: Input split: file:/opt/application/data.txt:562740360+19404840

22/12/19 11:52:02 INFO PythonRunner: Times: total = 8304, boot = -97, init = 183, finish = 8218

22/12/19 11:52:02 INFO Executor: Finished task 15.0 in stage 0.0 (TID 15). 2194 bytes result sent to driver

22/12/19 11:52:02 INFO TaskSetManager: Starting task 30.0 in stage 0.0 (TID 30) (d0a278a3161b, executor driver, partition 30, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:02 INFO TaskSetManager: Finished task 15.0 in stage 0.0 (TID 15) in 8380 ms on d0a278a3161b (executor driver) (19/100)

22/12/19 11:52:02 INFO Executor: Running task 30.0 in stage 0.0 (TID 30)

22/12/19 11:52:02 INFO HadoopRDD: Input split: file:/opt/application/data.txt:582145200+19404840

22/12/19 11:52:04 INFO PythonRunner: Times: total = 5683, boot = 7, init = 32, finish = 5644

22/12/19 11:52:04 INFO Executor: Finished task 20.0 in stage 0.0 (TID 20). 2195 bytes result sent to driver

22/12/19 11:52:04 INFO TaskSetManager: Starting task 31.0 in stage 0.0 (TID 31) (d0a278a3161b, executor driver, partition 31, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:04 INFO Executor: Running task 31.0 in stage 0.0 (TID 31)

22/12/19 11:52:04 INFO TaskSetManager: Finished task 20.0 in stage 0.0 (TID 20) in 5720 ms on d0a278a3161b (executor driver) (20/100)

22/12/19 11:52:04 INFO HadoopRDD: Input split: file:/opt/application/data.txt:601550040+19404840

22/12/19 11:52:04 INFO PythonRunner: Times: total = 9786, boot = -199, init = 218, finish = 9767

22/12/19 11:52:04 INFO Executor: Finished task 19.0 in stage 0.0 (TID 19). 2206 bytes result sent to driver

22/12/19 11:52:04 INFO TaskSetManager: Starting task 32.0 in stage 0.0 (TID 32) (d0a278a3161b, executor driver, partition 32, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:04 INFO TaskSetManager: Finished task 19.0 in stage 0.0 (TID 19) in 9919 ms on d0a278a3161b (executor driver) (21/100)

22/12/19 11:52:04 INFO Executor: Running task 32.0 in stage 0.0 (TID 32)

22/12/19 11:52:04 INFO HadoopRDD: Input split: file:/opt/application/data.txt:620954880+19404840

22/12/19 11:52:06 INFO PythonRunner: Times: total = 5915, boot = 273, init = 0, finish = 5642

22/12/19 11:52:06 INFO Executor: Finished task 22.0 in stage 0.0 (TID 22). 2217 bytes result sent to driver

22/12/19 11:52:06 INFO TaskSetManager: Starting task 33.0 in stage 0.0 (TID 33) (d0a278a3161b, executor driver, partition 33, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:06 INFO TaskSetManager: Finished task 22.0 in stage 0.0 (TID 22) in 5930 ms on d0a278a3161b (executor driver) (22/100)

22/12/19 11:52:06 INFO Executor: Running task 33.0 in stage 0.0 (TID 33)

22/12/19 11:52:06 INFO HadoopRDD: Input split: file:/opt/application/data.txt:640359720+19404840

22/12/19 11:52:06 INFO PythonRunner: Times: total = 4920, boot = -1, init = 3, finish = 4918

22/12/19 11:52:06 INFO Executor: Finished task 26.0 in stage 0.0 (TID 26). 2199 bytes result sent to driver

22/12/19 11:52:06 INFO TaskSetManager: Starting task 34.0 in stage 0.0 (TID 34) (d0a278a3161b, executor driver, partition 34, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:06 INFO TaskSetManager: Finished task 26.0 in stage 0.0 (TID 26) in 4943 ms on d0a278a3161b (executor driver) (23/100)

22/12/19 11:52:06 INFO Executor: Running task 34.0 in stage 0.0 (TID 34)

22/12/19 11:52:06 INFO HadoopRDD: Input split: file:/opt/application/data.txt:659764560+19404840

22/12/19 11:52:06 INFO PythonRunner: Times: total = 5913, boot = -17, init = 48, finish = 5882

22/12/19 11:52:06 INFO Executor: Finished task 24.0 in stage 0.0 (TID 24). 2197 bytes result sent to driver

22/12/19 11:52:06 INFO TaskSetManager: Starting task 35.0 in stage 0.0 (TID 35) (d0a278a3161b, executor driver, partition 35, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:06 INFO TaskSetManager: Finished task 24.0 in stage 0.0 (TID 24) in 5927 ms on d0a278a3161b (executor driver) (24/100)

22/12/19 11:52:06 INFO Executor: Running task 35.0 in stage 0.0 (TID 35)

22/12/19 11:52:06 INFO HadoopRDD: Input split: file:/opt/application/data.txt:679169400+19404840

22/12/19 11:52:07 INFO PythonRunner: Times: total = 5087, boot = -22, init = 25, finish = 5084

22/12/19 11:52:07 INFO Executor: Finished task 29.0 in stage 0.0 (TID 29). 2203 bytes result sent to driver

22/12/19 11:52:07 INFO TaskSetManager: Starting task 36.0 in stage 0.0 (TID 36) (d0a278a3161b, executor driver, partition 36, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:07 INFO TaskSetManager: Finished task 29.0 in stage 0.0 (TID 29) in 5105 ms on d0a278a3161b (executor driver) (25/100)

22/12/19 11:52:07 INFO Executor: Running task 36.0 in stage 0.0 (TID 36)

22/12/19 11:52:07 INFO HadoopRDD: Input split: file:/opt/application/data.txt:698574240+19404840

22/12/19 11:52:07 INFO PythonRunner: Times: total = 6725, boot = -177, init = 208, finish = 6694

22/12/19 11:52:07 INFO Executor: Finished task 25.0 in stage 0.0 (TID 25). 2204 bytes result sent to driver

22/12/19 11:52:07 INFO TaskSetManager: Starting task 37.0 in stage 0.0 (TID 37) (d0a278a3161b, executor driver, partition 37, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:07 INFO Executor: Running task 37.0 in stage 0.0 (TID 37)

22/12/19 11:52:07 INFO HadoopRDD: Input split: file:/opt/application/data.txt:717979080+19404840

22/12/19 11:52:07 INFO TaskSetManager: Finished task 25.0 in stage 0.0 (TID 25) in 6773 ms on d0a278a3161b (executor driver) (26/100)

22/12/19 11:52:07 INFO PythonRunner: Times: total = 5136, boot = -25, init = 41, finish = 5120

22/12/19 11:52:07 INFO Executor: Finished task 30.0 in stage 0.0 (TID 30). 2194 bytes result sent to driver

22/12/19 11:52:07 INFO TaskSetManager: Starting task 38.0 in stage 0.0 (TID 38) (d0a278a3161b, executor driver, partition 38, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:07 INFO TaskSetManager: Finished task 30.0 in stage 0.0 (TID 30) in 5159 ms on d0a278a3161b (executor driver) (27/100)

22/12/19 11:52:07 INFO Executor: Running task 38.0 in stage 0.0 (TID 38)

22/12/19 11:52:07 INFO HadoopRDD: Input split: file:/opt/application/data.txt:737383920+19404840

22/12/19 11:52:07 INFO PythonRunner: Times: total = 9215, boot = -35, init = 37, finish = 9213

22/12/19 11:52:07 INFO Executor: Finished task 21.0 in stage 0.0 (TID 21). 2189 bytes result sent to driver

22/12/19 11:52:07 INFO TaskSetManager: Starting task 39.0 in stage 0.0 (TID 39) (d0a278a3161b, executor driver, partition 39, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:07 INFO Executor: Running task 39.0 in stage 0.0 (TID 39)

22/12/19 11:52:07 INFO TaskSetManager: Finished task 21.0 in stage 0.0 (TID 21) in 9251 ms on d0a278a3161b (executor driver) (28/100)

22/12/19 11:52:07 INFO HadoopRDD: Input split: file:/opt/application/data.txt:756788760+19404840

22/12/19 11:52:08 INFO PythonRunner: Times: total = 5996, boot = -58, init = 66, finish = 5988

22/12/19 11:52:08 INFO Executor: Finished task 28.0 in stage 0.0 (TID 28). 2201 bytes result sent to driver

22/12/19 11:52:08 INFO TaskSetManager: Starting task 40.0 in stage 0.0 (TID 40) (d0a278a3161b, executor driver, partition 40, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:08 INFO Executor: Running task 40.0 in stage 0.0 (TID 40)

22/12/19 11:52:08 INFO TaskSetManager: Finished task 28.0 in stage 0.0 (TID 28) in 6024 ms on d0a278a3161b (executor driver) (29/100)

22/12/19 11:52:08 INFO HadoopRDD: Input split: file:/opt/application/data.txt:776193600+19404840

22/12/19 11:52:09 INFO PythonRunner: Times: total = 8860, boot = -4, init = 15, finish = 8849

22/12/19 11:52:09 INFO Executor: Finished task 23.0 in stage 0.0 (TID 23). 2183 bytes result sent to driver

22/12/19 11:52:09 INFO TaskSetManager: Starting task 41.0 in stage 0.0 (TID 41) (d0a278a3161b, executor driver, partition 41, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:09 INFO TaskSetManager: Finished task 23.0 in stage 0.0 (TID 23) in 8887 ms on d0a278a3161b (executor driver) (30/100)

22/12/19 11:52:09 INFO Executor: Running task 41.0 in stage 0.0 (TID 41)

22/12/19 11:52:09 INFO HadoopRDD: Input split: file:/opt/application/data.txt:795598440+19404840

22/12/19 11:52:10 INFO PythonRunner: Times: total = 5976, boot = -3, init = 18, finish = 5961

22/12/19 11:52:10 INFO Executor: Finished task 31.0 in stage 0.0 (TID 31). 2204 bytes result sent to driver

22/12/19 11:52:10 INFO TaskSetManager: Starting task 42.0 in stage 0.0 (TID 42) (d0a278a3161b, executor driver, partition 42, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:10 INFO Executor: Running task 42.0 in stage 0.0 (TID 42)

22/12/19 11:52:10 INFO TaskSetManager: Finished task 31.0 in stage 0.0 (TID 31) in 5996 ms on d0a278a3161b (executor driver) (31/100)

22/12/19 11:52:10 INFO HadoopRDD: Input split: file:/opt/application/data.txt:815003280+19404840

22/12/19 11:52:10 INFO PythonRunner: Times: total = 8234, boot = -35, init = 40, finish = 8229

22/12/19 11:52:10 INFO Executor: Finished task 27.0 in stage 0.0 (TID 27). 2215 bytes result sent to driver

22/12/19 11:52:10 INFO TaskSetManager: Starting task 43.0 in stage 0.0 (TID 43) (d0a278a3161b, executor driver, partition 43, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:10 INFO TaskSetManager: Finished task 27.0 in stage 0.0 (TID 27) in 8283 ms on d0a278a3161b (executor driver) (32/100)

22/12/19 11:52:10 INFO Executor: Running task 43.0 in stage 0.0 (TID 43)

22/12/19 11:52:10 INFO HadoopRDD: Input split: file:/opt/application/data.txt:834408120+19404840

22/12/19 11:52:11 INFO PythonRunner: Times: total = 5501, boot = -27, init = 38, finish = 5490

22/12/19 11:52:11 INFO Executor: Finished task 33.0 in stage 0.0 (TID 33). 2196 bytes result sent to driver

22/12/19 11:52:11 INFO TaskSetManager: Starting task 44.0 in stage 0.0 (TID 44) (d0a278a3161b, executor driver, partition 44, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:11 INFO TaskSetManager: Finished task 33.0 in stage 0.0 (TID 33) in 5543 ms on d0a278a3161b (executor driver) (33/100)

22/12/19 11:52:11 INFO Executor: Running task 44.0 in stage 0.0 (TID 44)

22/12/19 11:52:11 INFO HadoopRDD: Input split: file:/opt/application/data.txt:853812960+19404840

22/12/19 11:52:11 INFO PythonRunner: Times: total = 5227, boot = -19, init = 49, finish = 5197

22/12/19 11:52:11 INFO Executor: Finished task 35.0 in stage 0.0 (TID 35). 2212 bytes result sent to driver

22/12/19 11:52:11 INFO TaskSetManager: Starting task 45.0 in stage 0.0 (TID 45) (d0a278a3161b, executor driver, partition 45, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:11 INFO TaskSetManager: Finished task 35.0 in stage 0.0 (TID 35) in 5314 ms on d0a278a3161b (executor driver) (34/100)

22/12/19 11:52:11 INFO Executor: Running task 45.0 in stage 0.0 (TID 45)

22/12/19 11:52:11 INFO HadoopRDD: Input split: file:/opt/application/data.txt:873217800+19404840

22/12/19 11:52:12 INFO PythonRunner: Times: total = 7218, boot = -16, init = 38, finish = 7196

22/12/19 11:52:12 INFO Executor: Finished task 32.0 in stage 0.0 (TID 32). 2179 bytes result sent to driver

22/12/19 11:52:12 INFO TaskSetManager: Starting task 46.0 in stage 0.0 (TID 46) (d0a278a3161b, executor driver, partition 46, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:12 INFO Executor: Running task 46.0 in stage 0.0 (TID 46)

22/12/19 11:52:12 INFO TaskSetManager: Finished task 32.0 in stage 0.0 (TID 32) in 7256 ms on d0a278a3161b (executor driver) (35/100)

22/12/19 11:52:12 INFO HadoopRDD: Input split: file:/opt/application/data.txt:892622640+19404840

22/12/19 11:52:12 INFO PythonRunner: Times: total = 6383, boot = -24, init = 29, finish = 6378

22/12/19 11:52:12 INFO Executor: Finished task 34.0 in stage 0.0 (TID 34). 2202 bytes result sent to driver

22/12/19 11:52:12 INFO TaskSetManager: Starting task 47.0 in stage 0.0 (TID 47) (d0a278a3161b, executor driver, partition 47, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:12 INFO TaskSetManager: Finished task 34.0 in stage 0.0 (TID 34) in 6412 ms on d0a278a3161b (executor driver) (36/100)

22/12/19 11:52:12 INFO Executor: Running task 47.0 in stage 0.0 (TID 47)

22/12/19 11:52:12 INFO HadoopRDD: Input split: file:/opt/application/data.txt:912027480+19404840

22/12/19 11:52:12 INFO PythonRunner: Times: total = 5107, boot = -27, init = 37, finish = 5097

22/12/19 11:52:12 INFO Executor: Finished task 37.0 in stage 0.0 (TID 37). 2187 bytes result sent to driver

22/12/19 11:52:12 INFO TaskSetManager: Starting task 48.0 in stage 0.0 (TID 48) (d0a278a3161b, executor driver, partition 48, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:12 INFO Executor: Running task 48.0 in stage 0.0 (TID 48)

22/12/19 11:52:12 INFO TaskSetManager: Finished task 37.0 in stage 0.0 (TID 37) in 5133 ms on d0a278a3161b (executor driver) (37/100)

22/12/19 11:52:12 INFO HadoopRDD: Input split: file:/opt/application/data.txt:931432320+19404840

22/12/19 11:52:13 INFO PythonRunner: Times: total = 5985, boot = -14, init = 38, finish = 5961

22/12/19 11:52:13 INFO Executor: Finished task 38.0 in stage 0.0 (TID 38). 2210 bytes result sent to driver

22/12/19 11:52:13 INFO TaskSetManager: Starting task 49.0 in stage 0.0 (TID 49) (d0a278a3161b, executor driver, partition 49, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:13 INFO TaskSetManager: Finished task 38.0 in stage 0.0 (TID 38) in 6007 ms on d0a278a3161b (executor driver) (38/100)

22/12/19 11:52:13 INFO Executor: Running task 49.0 in stage 0.0 (TID 49)

22/12/19 11:52:13 INFO HadoopRDD: Input split: file:/opt/application/data.txt:950837160+19404840

22/12/19 11:52:14 INFO PythonRunner: Times: total = 6371, boot = 2, init = 12, finish = 6357

22/12/19 11:52:14 INFO Executor: Finished task 40.0 in stage 0.0 (TID 40). 2195 bytes result sent to driver

22/12/19 11:52:14 INFO TaskSetManager: Starting task 50.0 in stage 0.0 (TID 50) (d0a278a3161b, executor driver, partition 50, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:14 INFO TaskSetManager: Finished task 40.0 in stage 0.0 (TID 40) in 6381 ms on d0a278a3161b (executor driver) (39/100)

22/12/19 11:52:14 INFO Executor: Running task 50.0 in stage 0.0 (TID 50)

22/12/19 11:52:14 INFO HadoopRDD: Input split: file:/opt/application/data.txt:970242000+19404840

22/12/19 11:52:16 INFO PythonRunner: Times: total = 6533, boot = -17, init = 35, finish = 6515

22/12/19 11:52:16 INFO PythonRunner: Times: total = 8098, boot = -18, init = 27, finish = 8089

22/12/19 11:52:16 INFO Executor: Finished task 41.0 in stage 0.0 (TID 41). 2203 bytes result sent to driver

22/12/19 11:52:16 INFO PythonRunner: Times: total = 7487, boot = -17, init = 39, finish = 7465

22/12/19 11:52:16 INFO PythonRunner: Times: total = 5144, boot = -2, init = 17, finish = 5129

22/12/19 11:52:16 INFO TaskSetManager: Starting task 51.0 in stage 0.0 (TID 51) (d0a278a3161b, executor driver, partition 51, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:16 INFO Executor: Finished task 36.0 in stage 0.0 (TID 36). 2221 bytes result sent to driver

22/12/19 11:52:17 INFO PythonRunner: Times: total = 5767, boot = 4, init = 21, finish = 5742

22/12/19 11:52:17 INFO Executor: Running task 51.0 in stage 0.0 (TID 51)

22/12/19 11:52:17 INFO TaskSetManager: Starting task 52.0 in stage 0.0 (TID 52) (d0a278a3161b, executor driver, partition 52, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:17 INFO TaskSetManager: Finished task 41.0 in stage 0.0 (TID 41) in 8471 ms on d0a278a3161b (executor driver) (40/100)

22/12/19 11:52:17 INFO TaskSetManager: Finished task 36.0 in stage 0.0 (TID 36) in 10281 ms on d0a278a3161b (executor driver) (41/100)

22/12/19 11:52:17 INFO HadoopRDD: Input split: file:/opt/application/data.txt:989646840+19404840

22/12/19 11:52:17 INFO Executor: Running task 52.0 in stage 0.0 (TID 52)

22/12/19 11:52:17 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1009051680+19404840

22/12/19 11:52:18 INFO PythonRunner: Times: total = 5175, boot = -12, init = 24, finish = 5163

22/12/19 11:52:18 INFO Executor: Finished task 44.0 in stage 0.0 (TID 44). 2195 bytes result sent to driver

22/12/19 11:52:18 INFO TaskSetManager: Starting task 53.0 in stage 0.0 (TID 53) (d0a278a3161b, executor driver, partition 53, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:18 INFO TaskSetManager: Finished task 44.0 in stage 0.0 (TID 44) in 7258 ms on d0a278a3161b (executor driver) (42/100)

22/12/19 11:52:20 INFO PythonRunner: Times: total = 7078, boot = -49, init = 84, finish = 7043

22/12/19 11:52:20 INFO AsyncEventQueue: Process of event SparkListenerTaskStart(0,0,org.apache.spark.scheduler.TaskInfo@7d329fe3) by listener AppStatusListener took 1.092034031s.

22/12/19 11:52:21 INFO Executor: Running task 53.0 in stage 0.0 (TID 53)

22/12/19 11:52:21 INFO PythonRunner: Times: total = 6821, boot = 0, init = 20, finish = 6801

22/12/19 11:52:21 INFO PythonRunner: Times: total = 8492, boot = 1, init = 5, finish = 8486

22/12/19 11:52:21 INFO PythonRunner: Times: total = 7799, boot = 2, init = 2, finish = 7795

22/12/19 11:52:21 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1028456520+19404840

22/12/19 11:52:21 INFO Executor: Finished task 42.0 in stage 0.0 (TID 42). 2216 bytes result sent to driver

22/12/19 11:52:21 INFO Executor: Finished task 45.0 in stage 0.0 (TID 45). 2189 bytes result sent to driver

22/12/19 11:52:22 INFO Executor: Finished task 47.0 in stage 0.0 (TID 47). 2191 bytes result sent to driver

22/12/19 11:52:22 INFO TaskSetManager: Starting task 54.0 in stage 0.0 (TID 54) (d0a278a3161b, executor driver, partition 54, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:21 INFO Executor: Finished task 39.0 in stage 0.0 (TID 39). 2198 bytes result sent to driver

22/12/19 11:52:22 INFO Executor: Finished task 43.0 in stage 0.0 (TID 43). 2199 bytes result sent to driver

22/12/19 11:52:22 INFO Executor: Finished task 46.0 in stage 0.0 (TID 46). 2209 bytes result sent to driver

22/12/19 11:52:22 INFO Executor: Finished task 48.0 in stage 0.0 (TID 48). 2186 bytes result sent to driver

22/12/19 11:52:22 INFO TaskSetManager: Finished task 42.0 in stage 0.0 (TID 42) in 12062 ms on d0a278a3161b (executor driver) (43/100)

22/12/19 11:52:22 INFO Executor: Running task 54.0 in stage 0.0 (TID 54)

22/12/19 11:52:22 INFO TaskSetManager: Starting task 55.0 in stage 0.0 (TID 55) (d0a278a3161b, executor driver, partition 55, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:22 INFO Executor: Running task 55.0 in stage 0.0 (TID 55)

22/12/19 11:52:22 INFO TaskSetManager: Starting task 56.0 in stage 0.0 (TID 56) (d0a278a3161b, executor driver, partition 56, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:22 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1047861360+19404840

22/12/19 11:52:22 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1067266200+19404840

22/12/19 11:52:22 INFO TaskSetManager: Starting task 57.0 in stage 0.0 (TID 57) (d0a278a3161b, executor driver, partition 57, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:22 INFO Executor: Running task 56.0 in stage 0.0 (TID 56)

22/12/19 11:52:22 INFO Executor: Running task 57.0 in stage 0.0 (TID 57)

22/12/19 11:52:22 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1106075880+19404840

22/12/19 11:52:22 INFO TaskSetManager: Starting task 58.0 in stage 0.0 (TID 58) (d0a278a3161b, executor driver, partition 58, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:22 INFO TaskSetManager: Finished task 45.0 in stage 0.0 (TID 45) in 10651 ms on d0a278a3161b (executor driver) (44/100)

22/12/19 11:52:22 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1086671040+19404840

22/12/19 11:52:22 INFO Executor: Running task 58.0 in stage 0.0 (TID 58)

22/12/19 11:52:22 INFO TaskSetManager: Finished task 47.0 in stage 0.0 (TID 47) in 9756 ms on d0a278a3161b (executor driver) (45/100)

22/12/19 11:52:22 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1125480720+19404840

22/12/19 11:52:22 INFO TaskSetManager: Finished task 39.0 in stage 0.0 (TID 39) in 14448 ms on d0a278a3161b (executor driver) (46/100)

22/12/19 11:52:22 INFO TaskSetManager: Starting task 59.0 in stage 0.0 (TID 59) (d0a278a3161b, executor driver, partition 59, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:22 INFO Executor: Running task 59.0 in stage 0.0 (TID 59)

22/12/19 11:52:22 INFO TaskSetManager: Starting task 60.0 in stage 0.0 (TID 60) (d0a278a3161b, executor driver, partition 60, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:22 INFO TaskSetManager: Finished task 43.0 in stage 0.0 (TID 43) in 12145 ms on d0a278a3161b (executor driver) (47/100)

22/12/19 11:52:22 INFO Executor: Running task 60.0 in stage 0.0 (TID 60)

22/12/19 11:52:22 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1164290400+19404840

22/12/19 11:52:22 INFO TaskSetManager: Finished task 46.0 in stage 0.0 (TID 46) in 10474 ms on d0a278a3161b (executor driver) (48/100)

22/12/19 11:52:22 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1144885560+19404840

22/12/19 11:52:22 INFO TaskSetManager: Finished task 48.0 in stage 0.0 (TID 48) in 9704 ms on d0a278a3161b (executor driver) (49/100)

22/12/19 11:52:22 INFO PythonRunner: Times: total = 8980, boot = 16, init = 2, finish = 8962

22/12/19 11:52:22 INFO Executor: Finished task 49.0 in stage 0.0 (TID 49). 2202 bytes result sent to driver

22/12/19 11:52:22 INFO TaskSetManager: Starting task 61.0 in stage 0.0 (TID 61) (d0a278a3161b, executor driver, partition 61, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:22 INFO Executor: Running task 61.0 in stage 0.0 (TID 61)

22/12/19 11:52:22 INFO TaskSetManager: Finished task 49.0 in stage 0.0 (TID 49) in 9068 ms on d0a278a3161b (executor driver) (50/100)

22/12/19 11:52:22 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1183695240+19404840

22/12/19 11:52:23 INFO PythonRunner: Times: total = 9071, boot = 6, init = 10, finish = 9055

22/12/19 11:52:23 INFO Executor: Finished task 50.0 in stage 0.0 (TID 50). 2207 bytes result sent to driver

22/12/19 11:52:23 INFO TaskSetManager: Starting task 62.0 in stage 0.0 (TID 62) (d0a278a3161b, executor driver, partition 62, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:23 INFO TaskSetManager: Finished task 50.0 in stage 0.0 (TID 50) in 9093 ms on d0a278a3161b (executor driver) (51/100)

22/12/19 11:52:23 INFO Executor: Running task 62.0 in stage 0.0 (TID 62)

22/12/19 11:52:23 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1203100080+19404840

22/12/19 11:52:27 INFO PythonRunner: Times: total = 5012, boot = -4441, init = 4493, finish = 4960

22/12/19 11:52:27 INFO Executor: Finished task 57.0 in stage 0.0 (TID 57). 2198 bytes result sent to driver

22/12/19 11:52:27 INFO TaskSetManager: Starting task 63.0 in stage 0.0 (TID 63) (d0a278a3161b, executor driver, partition 63, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:27 INFO TaskSetManager: Finished task 57.0 in stage 0.0 (TID 57) in 5027 ms on d0a278a3161b (executor driver) (52/100)

22/12/19 11:52:27 INFO Executor: Running task 63.0 in stage 0.0 (TID 63)

22/12/19 11:52:27 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1222504920+19404840

22/12/19 11:52:27 INFO PythonRunner: Times: total = 5077, boot = -3570, init = 3616, finish = 5031

22/12/19 11:52:27 INFO Executor: Finished task 58.0 in stage 0.0 (TID 58). 2194 bytes result sent to driver

22/12/19 11:52:27 INFO TaskSetManager: Starting task 64.0 in stage 0.0 (TID 64) (d0a278a3161b, executor driver, partition 64, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:27 INFO TaskSetManager: Finished task 58.0 in stage 0.0 (TID 58) in 5107 ms on d0a278a3161b (executor driver) (53/100)

22/12/19 11:52:27 INFO Executor: Running task 64.0 in stage 0.0 (TID 64)

22/12/19 11:52:27 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1241909760+19404840

22/12/19 11:52:27 INFO PythonRunner: Times: total = 5299, boot = -2026, init = 2074, finish = 5251

22/12/19 11:52:27 INFO Executor: Finished task 60.0 in stage 0.0 (TID 60). 2194 bytes result sent to driver

22/12/19 11:52:27 INFO TaskSetManager: Starting task 65.0 in stage 0.0 (TID 65) (d0a278a3161b, executor driver, partition 65, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:27 INFO Executor: Running task 65.0 in stage 0.0 (TID 65)

22/12/19 11:52:27 INFO TaskSetManager: Finished task 60.0 in stage 0.0 (TID 60) in 5334 ms on d0a278a3161b (executor driver) (54/100)

22/12/19 11:52:27 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1261314600+19404840

22/12/19 11:52:28 INFO PythonRunner: Times: total = 7295, boot = -4364, init = 4379, finish = 7280

22/12/19 11:52:28 INFO Executor: Finished task 53.0 in stage 0.0 (TID 53). 2195 bytes result sent to driver

22/12/19 11:52:28 INFO TaskSetManager: Starting task 66.0 in stage 0.0 (TID 66) (d0a278a3161b, executor driver, partition 66, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:28 INFO Executor: Running task 66.0 in stage 0.0 (TID 66)

22/12/19 11:52:28 INFO TaskSetManager: Finished task 53.0 in stage 0.0 (TID 53) in 9567 ms on d0a278a3161b (executor driver) (55/100)

22/12/19 11:52:28 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1280719440+19404840

22/12/19 11:52:28 INFO PythonRunner: Times: total = 4855, boot = -7, init = 18, finish = 4844

22/12/19 11:52:28 INFO Executor: Finished task 62.0 in stage 0.0 (TID 62). 2191 bytes result sent to driver

22/12/19 11:52:28 INFO TaskSetManager: Starting task 67.0 in stage 0.0 (TID 67) (d0a278a3161b, executor driver, partition 67, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:28 INFO Executor: Running task 67.0 in stage 0.0 (TID 67)

22/12/19 11:52:28 INFO TaskSetManager: Finished task 62.0 in stage 0.0 (TID 62) in 4868 ms on d0a278a3161b (executor driver) (56/100)

22/12/19 11:52:28 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1300124280+19404840

22/12/19 11:52:28 INFO PythonRunner: Times: total = 10894, boot = -1911, init = 4559, finish = 8246

22/12/19 11:52:28 INFO Executor: Finished task 51.0 in stage 0.0 (TID 51). 2205 bytes result sent to driver

22/12/19 11:52:28 INFO TaskSetManager: Starting task 68.0 in stage 0.0 (TID 68) (d0a278a3161b, executor driver, partition 68, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:28 INFO TaskSetManager: Finished task 51.0 in stage 0.0 (TID 51) in 12119 ms on d0a278a3161b (executor driver) (57/100)

22/12/19 11:52:28 INFO Executor: Running task 68.0 in stage 0.0 (TID 68)

22/12/19 11:52:28 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1319529120+19404840

22/12/19 11:52:28 INFO PythonRunner: Times: total = 6397, boot = -6304, init = 6381, finish = 6320

22/12/19 11:52:28 INFO Executor: Finished task 54.0 in stage 0.0 (TID 54). 2201 bytes result sent to driver

22/12/19 11:52:28 INFO TaskSetManager: Starting task 69.0 in stage 0.0 (TID 69) (d0a278a3161b, executor driver, partition 69, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:28 INFO Executor: Running task 69.0 in stage 0.0 (TID 69)

22/12/19 11:52:28 INFO TaskSetManager: Finished task 54.0 in stage 0.0 (TID 54) in 6430 ms on d0a278a3161b (executor driver) (58/100)

22/12/19 11:52:28 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1338933960+19404840

22/12/19 11:52:29 INFO PythonRunner: Times: total = 6368, boot = -82, init = 105, finish = 6345

22/12/19 11:52:29 INFO Executor: Finished task 61.0 in stage 0.0 (TID 61). 2208 bytes result sent to driver

22/12/19 11:52:29 INFO TaskSetManager: Starting task 70.0 in stage 0.0 (TID 70) (d0a278a3161b, executor driver, partition 70, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:29 INFO TaskSetManager: Finished task 61.0 in stage 0.0 (TID 61) in 6404 ms on d0a278a3161b (executor driver) (59/100)

22/12/19 11:52:29 INFO Executor: Running task 70.0 in stage 0.0 (TID 70)

22/12/19 11:52:29 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1358338800+19404840

22/12/19 11:52:29 INFO PythonRunner: Times: total = 11761, boot = -2169, init = 4807, finish = 9123

22/12/19 11:52:29 INFO Executor: Finished task 52.0 in stage 0.0 (TID 52). 2188 bytes result sent to driver

22/12/19 11:52:29 INFO TaskSetManager: Starting task 71.0 in stage 0.0 (TID 71) (d0a278a3161b, executor driver, partition 71, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:29 INFO Executor: Running task 71.0 in stage 0.0 (TID 71)

22/12/19 11:52:29 INFO TaskSetManager: Finished task 52.0 in stage 0.0 (TID 52) in 11779 ms on d0a278a3161b (executor driver) (60/100)

22/12/19 11:52:29 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1377743640+19404840

22/12/19 11:52:29 INFO PythonRunner: Times: total = 7539, boot = -6931, init = 6996, finish = 7474

22/12/19 11:52:29 INFO Executor: Finished task 55.0 in stage 0.0 (TID 55). 2206 bytes result sent to driver

22/12/19 11:52:29 INFO TaskSetManager: Starting task 72.0 in stage 0.0 (TID 72) (d0a278a3161b, executor driver, partition 72, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:29 INFO Executor: Running task 72.0 in stage 0.0 (TID 72)

22/12/19 11:52:29 INFO TaskSetManager: Finished task 55.0 in stage 0.0 (TID 55) in 7557 ms on d0a278a3161b (executor driver) (61/100)

22/12/19 11:52:29 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1397148480+19404840

22/12/19 11:52:30 INFO PythonRunner: Times: total = 8048, boot = -3567, init = 3627, finish = 7988

22/12/19 11:52:30 INFO Executor: Finished task 56.0 in stage 0.0 (TID 56). 2207 bytes result sent to driver

22/12/19 11:52:30 INFO TaskSetManager: Starting task 73.0 in stage 0.0 (TID 73) (d0a278a3161b, executor driver, partition 73, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:30 INFO TaskSetManager: Finished task 56.0 in stage 0.0 (TID 56) in 8069 ms on d0a278a3161b (executor driver) (62/100)

22/12/19 11:52:30 INFO Executor: Running task 73.0 in stage 0.0 (TID 73)

22/12/19 11:52:30 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1416553320+19404840

22/12/19 11:52:30 INFO PythonRunner: Times: total = 8008, boot = -3643, init = 3673, finish = 7978

22/12/19 11:52:30 INFO Executor: Finished task 59.0 in stage 0.0 (TID 59). 2199 bytes result sent to driver

22/12/19 11:52:30 INFO TaskSetManager: Starting task 74.0 in stage 0.0 (TID 74) (d0a278a3161b, executor driver, partition 74, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:30 INFO TaskSetManager: Finished task 59.0 in stage 0.0 (TID 59) in 8083 ms on d0a278a3161b (executor driver) (63/100)

22/12/19 11:52:30 INFO Executor: Running task 74.0 in stage 0.0 (TID 74)

22/12/19 11:52:30 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1435958160+19404840

22/12/19 11:52:32 INFO PythonRunner: Times: total = 5036, boot = -14, init = 39, finish = 5011

22/12/19 11:52:32 INFO Executor: Finished task 64.0 in stage 0.0 (TID 64). 2193 bytes result sent to driver

22/12/19 11:52:32 INFO TaskSetManager: Starting task 75.0 in stage 0.0 (TID 75) (d0a278a3161b, executor driver, partition 75, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:32 INFO TaskSetManager: Finished task 64.0 in stage 0.0 (TID 64) in 5051 ms on d0a278a3161b (executor driver) (64/100)

22/12/19 11:52:32 INFO Executor: Running task 75.0 in stage 0.0 (TID 75)

22/12/19 11:52:32 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1455363000+19404840

22/12/19 11:52:33 INFO PythonRunner: Times: total = 4795, boot = 6, init = 34, finish = 4755

22/12/19 11:52:33 INFO Executor: Finished task 68.0 in stage 0.0 (TID 68). 2204 bytes result sent to driver

22/12/19 11:52:33 INFO TaskSetManager: Starting task 76.0 in stage 0.0 (TID 76) (d0a278a3161b, executor driver, partition 76, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:33 INFO Executor: Running task 76.0 in stage 0.0 (TID 76)

22/12/19 11:52:33 INFO TaskSetManager: Finished task 68.0 in stage 0.0 (TID 68) in 4806 ms on d0a278a3161b (executor driver) (65/100)

22/12/19 11:52:33 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1474767840+19404840

22/12/19 11:52:33 INFO PythonRunner: Times: total = 4918, boot = -14, init = 16, finish = 4916

22/12/19 11:52:33 INFO Executor: Finished task 67.0 in stage 0.0 (TID 67). 2191 bytes result sent to driver

22/12/19 11:52:33 INFO TaskSetManager: Starting task 77.0 in stage 0.0 (TID 77) (d0a278a3161b, executor driver, partition 77, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:33 INFO Executor: Running task 77.0 in stage 0.0 (TID 77)

22/12/19 11:52:33 INFO TaskSetManager: Finished task 67.0 in stage 0.0 (TID 67) in 4948 ms on d0a278a3161b (executor driver) (66/100)

22/12/19 11:52:33 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1494172680+19404840

22/12/19 11:52:34 INFO PythonRunner: Times: total = 6578, boot = -2, init = 15, finish = 6565

22/12/19 11:52:34 INFO Executor: Finished task 63.0 in stage 0.0 (TID 63). 2207 bytes result sent to driver

22/12/19 11:52:34 INFO TaskSetManager: Starting task 78.0 in stage 0.0 (TID 78) (d0a278a3161b, executor driver, partition 78, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:34 INFO TaskSetManager: Finished task 63.0 in stage 0.0 (TID 63) in 6590 ms on d0a278a3161b (executor driver) (67/100)

22/12/19 11:52:34 INFO Executor: Running task 78.0 in stage 0.0 (TID 78)

22/12/19 11:52:34 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1513577520+19404840

22/12/19 11:52:34 INFO PythonRunner: Times: total = 4868, boot = -28, init = 38, finish = 4858

22/12/19 11:52:34 INFO Executor: Finished task 70.0 in stage 0.0 (TID 70). 2191 bytes result sent to driver

22/12/19 11:52:34 INFO TaskSetManager: Starting task 79.0 in stage 0.0 (TID 79) (d0a278a3161b, executor driver, partition 79, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:34 INFO TaskSetManager: Finished task 70.0 in stage 0.0 (TID 70) in 4889 ms on d0a278a3161b (executor driver) (68/100)

22/12/19 11:52:34 INFO Executor: Running task 79.0 in stage 0.0 (TID 79)

22/12/19 11:52:34 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1532982360+19404840

22/12/19 11:52:35 INFO PythonRunner: Times: total = 6342, boot = 6, init = 8, finish = 6328

22/12/19 11:52:35 INFO Executor: Finished task 69.0 in stage 0.0 (TID 69). 2198 bytes result sent to driver

22/12/19 11:52:35 INFO TaskSetManager: Starting task 80.0 in stage 0.0 (TID 80) (d0a278a3161b, executor driver, partition 80, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:35 INFO Executor: Running task 80.0 in stage 0.0 (TID 80)

22/12/19 11:52:35 INFO TaskSetManager: Finished task 69.0 in stage 0.0 (TID 69) in 6352 ms on d0a278a3161b (executor driver) (69/100)

22/12/19 11:52:35 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1552387200+19404840

22/12/19 11:52:35 INFO PythonRunner: Times: total = 5954, boot = 3, init = 13, finish = 5938

22/12/19 11:52:35 INFO Executor: Finished task 71.0 in stage 0.0 (TID 71). 2188 bytes result sent to driver

22/12/19 11:52:35 INFO TaskSetManager: Starting task 81.0 in stage 0.0 (TID 81) (d0a278a3161b, executor driver, partition 81, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:35 INFO Executor: Running task 81.0 in stage 0.0 (TID 81)

22/12/19 11:52:35 INFO TaskSetManager: Finished task 71.0 in stage 0.0 (TID 71) in 5964 ms on d0a278a3161b (executor driver) (70/100)

22/12/19 11:52:35 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1571792040+19404840

22/12/19 11:52:35 INFO PythonRunner: Times: total = 7170, boot = 7, init = 19, finish = 7144

22/12/19 11:52:35 INFO Executor: Finished task 66.0 in stage 0.0 (TID 66). 2199 bytes result sent to driver

22/12/19 11:52:35 INFO TaskSetManager: Starting task 82.0 in stage 0.0 (TID 82) (d0a278a3161b, executor driver, partition 82, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:35 INFO Executor: Running task 82.0 in stage 0.0 (TID 82)

22/12/19 11:52:35 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1591196880+19404840

22/12/19 11:52:35 INFO TaskSetManager: Finished task 66.0 in stage 0.0 (TID 66) in 7180 ms on d0a278a3161b (executor driver) (71/100)

22/12/19 11:52:36 INFO PythonRunner: Times: total = 8468, boot = 7, init = 13, finish = 8448

22/12/19 11:52:36 INFO Executor: Finished task 65.0 in stage 0.0 (TID 65). 2206 bytes result sent to driver

22/12/19 11:52:36 INFO TaskSetManager: Starting task 83.0 in stage 0.0 (TID 83) (d0a278a3161b, executor driver, partition 83, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:36 INFO TaskSetManager: Finished task 65.0 in stage 0.0 (TID 65) in 8478 ms on d0a278a3161b (executor driver) (72/100)

22/12/19 11:52:36 INFO Executor: Running task 83.0 in stage 0.0 (TID 83)

22/12/19 11:52:36 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1610601720+19404840

22/12/19 11:52:37 INFO PythonRunner: Times: total = 7514, boot = -2, init = 31, finish = 7485

22/12/19 11:52:37 INFO Executor: Finished task 72.0 in stage 0.0 (TID 72). 2200 bytes result sent to driver

22/12/19 11:52:37 INFO TaskSetManager: Starting task 84.0 in stage 0.0 (TID 84) (d0a278a3161b, executor driver, partition 84, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:37 INFO Executor: Running task 84.0 in stage 0.0 (TID 84)

22/12/19 11:52:37 INFO TaskSetManager: Finished task 72.0 in stage 0.0 (TID 72) in 7523 ms on d0a278a3161b (executor driver) (73/100)

22/12/19 11:52:37 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1630006560+19404840

22/12/19 11:52:38 INFO PythonRunner: Times: total = 5652, boot = 5, init = 24, finish = 5623

22/12/19 11:52:38 INFO Executor: Finished task 75.0 in stage 0.0 (TID 75). 2212 bytes result sent to driver

22/12/19 11:52:38 INFO TaskSetManager: Starting task 85.0 in stage 0.0 (TID 85) (d0a278a3161b, executor driver, partition 85, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:38 INFO TaskSetManager: Finished task 75.0 in stage 0.0 (TID 75) in 5659 ms on d0a278a3161b (executor driver) (74/100)

22/12/19 11:52:38 INFO Executor: Running task 85.0 in stage 0.0 (TID 85)

22/12/19 11:52:38 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1649411400+19404840

22/12/19 11:52:38 INFO PythonRunner: Times: total = 7943, boot = 2, init = 20, finish = 7921

22/12/19 11:52:38 INFO Executor: Finished task 74.0 in stage 0.0 (TID 74). 2198 bytes result sent to driver

22/12/19 11:52:38 INFO TaskSetManager: Starting task 86.0 in stage 0.0 (TID 86) (d0a278a3161b, executor driver, partition 86, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:38 INFO Executor: Running task 86.0 in stage 0.0 (TID 86)

22/12/19 11:52:38 INFO TaskSetManager: Finished task 74.0 in stage 0.0 (TID 74) in 7959 ms on d0a278a3161b (executor driver) (75/100)

22/12/19 11:52:38 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1668816240+19404840

22/12/19 11:52:39 INFO PythonRunner: Times: total = 8842, boot = 19, init = 0, finish = 8823

22/12/19 11:52:39 INFO Executor: Finished task 73.0 in stage 0.0 (TID 73). 2190 bytes result sent to driver

22/12/19 11:52:39 INFO TaskSetManager: Starting task 87.0 in stage 0.0 (TID 87) (d0a278a3161b, executor driver, partition 87, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:39 INFO TaskSetManager: Finished task 73.0 in stage 0.0 (TID 73) in 8879 ms on d0a278a3161b (executor driver) (76/100)

22/12/19 11:52:39 INFO Executor: Running task 87.0 in stage 0.0 (TID 87)

22/12/19 11:52:39 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1688221080+19404840

22/12/19 11:52:40 INFO PythonRunner: Times: total = 6775, boot = 5, init = 12, finish = 6758

22/12/19 11:52:40 INFO Executor: Finished task 76.0 in stage 0.0 (TID 76). 2191 bytes result sent to driver

22/12/19 11:52:40 INFO TaskSetManager: Starting task 88.0 in stage 0.0 (TID 88) (d0a278a3161b, executor driver, partition 88, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:40 INFO TaskSetManager: Finished task 76.0 in stage 0.0 (TID 76) in 6784 ms on d0a278a3161b (executor driver) (77/100)

22/12/19 11:52:40 INFO Executor: Running task 88.0 in stage 0.0 (TID 88)

22/12/19 11:52:40 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1707625920+19404840

22/12/19 11:52:40 INFO PythonRunner: Times: total = 7328, boot = 4, init = 13, finish = 7311

22/12/19 11:52:40 INFO Executor: Finished task 77.0 in stage 0.0 (TID 77). 2201 bytes result sent to driver

22/12/19 11:52:40 INFO TaskSetManager: Starting task 89.0 in stage 0.0 (TID 89) (d0a278a3161b, executor driver, partition 89, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:40 INFO Executor: Running task 89.0 in stage 0.0 (TID 89)

22/12/19 11:52:40 INFO TaskSetManager: Finished task 77.0 in stage 0.0 (TID 77) in 7340 ms on d0a278a3161b (executor driver) (78/100)

22/12/19 11:52:40 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1727030760+19404840

22/12/19 11:52:40 INFO PythonRunner: Times: total = 6791, boot = 7, init = 8, finish = 6776

22/12/19 11:52:40 INFO Executor: Finished task 79.0 in stage 0.0 (TID 79). 2205 bytes result sent to driver

22/12/19 11:52:40 INFO TaskSetManager: Starting task 90.0 in stage 0.0 (TID 90) (d0a278a3161b, executor driver, partition 90, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:40 INFO Executor: Running task 90.0 in stage 0.0 (TID 90)

22/12/19 11:52:40 INFO TaskSetManager: Finished task 79.0 in stage 0.0 (TID 79) in 6845 ms on d0a278a3161b (executor driver) (79/100)

22/12/19 11:52:40 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1746435600+19404840

22/12/19 11:52:40 INFO PythonRunner: Times: total = 6973, boot = 4, init = 1, finish = 6968

22/12/19 11:52:40 INFO Executor: Finished task 78.0 in stage 0.0 (TID 78). 2206 bytes result sent to driver

22/12/19 11:52:40 INFO TaskSetManager: Starting task 91.0 in stage 0.0 (TID 91) (d0a278a3161b, executor driver, partition 91, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:40 INFO Executor: Running task 91.0 in stage 0.0 (TID 91)

22/12/19 11:52:40 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1765840440+19404840

22/12/19 11:52:41 INFO TaskSetManager: Finished task 78.0 in stage 0.0 (TID 78) in 6990 ms on d0a278a3161b (executor driver) (80/100)

22/12/19 11:52:41 INFO PythonRunner: Times: total = 6778, boot = 22, init = 0, finish = 6756

22/12/19 11:52:41 INFO Executor: Finished task 80.0 in stage 0.0 (TID 80). 2210 bytes result sent to driver

22/12/19 11:52:41 INFO TaskSetManager: Starting task 92.0 in stage 0.0 (TID 92) (d0a278a3161b, executor driver, partition 92, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:41 INFO TaskSetManager: Finished task 80.0 in stage 0.0 (TID 80) in 6787 ms on d0a278a3161b (executor driver) (81/100)

22/12/19 11:52:41 INFO Executor: Running task 92.0 in stage 0.0 (TID 92)

22/12/19 11:52:41 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1785245280+19404840

22/12/19 11:52:42 INFO PythonRunner: Times: total = 6548, boot = 6, init = 9, finish = 6533

22/12/19 11:52:42 INFO Executor: Finished task 82.0 in stage 0.0 (TID 82). 2216 bytes result sent to driver

22/12/19 11:52:42 INFO TaskSetManager: Starting task 93.0 in stage 0.0 (TID 93) (d0a278a3161b, executor driver, partition 93, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:42 INFO Executor: Running task 93.0 in stage 0.0 (TID 93)

22/12/19 11:52:42 INFO TaskSetManager: Finished task 82.0 in stage 0.0 (TID 82) in 6558 ms on d0a278a3161b (executor driver) (82/100)

22/12/19 11:52:42 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1804650120+19404840

22/12/19 11:52:42 INFO PythonRunner: Times: total = 7064, boot = 4, init = 18, finish = 7042

22/12/19 11:52:42 INFO Executor: Finished task 81.0 in stage 0.0 (TID 81). 2208 bytes result sent to driver

22/12/19 11:52:42 INFO TaskSetManager: Starting task 94.0 in stage 0.0 (TID 94) (d0a278a3161b, executor driver, partition 94, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:42 INFO TaskSetManager: Finished task 81.0 in stage 0.0 (TID 81) in 7075 ms on d0a278a3161b (executor driver) (83/100)

22/12/19 11:52:42 INFO Executor: Running task 94.0 in stage 0.0 (TID 94)

22/12/19 11:52:42 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1824054960+19404840

22/12/19 11:52:43 INFO PythonRunner: Times: total = 6780, boot = 7, init = 10, finish = 6763

22/12/19 11:52:43 INFO Executor: Finished task 83.0 in stage 0.0 (TID 83). 2189 bytes result sent to driver

22/12/19 11:52:43 INFO TaskSetManager: Starting task 95.0 in stage 0.0 (TID 95) (d0a278a3161b, executor driver, partition 95, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:43 INFO Executor: Running task 95.0 in stage 0.0 (TID 95)

22/12/19 11:52:43 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1843459800+19404840

22/12/19 11:52:43 INFO TaskSetManager: Finished task 83.0 in stage 0.0 (TID 83) in 6795 ms on d0a278a3161b (executor driver) (84/100)

22/12/19 11:52:43 INFO PythonRunner: Times: total = 5644, boot = 9, init = 1, finish = 5634

22/12/19 11:52:43 INFO Executor: Finished task 85.0 in stage 0.0 (TID 85). 2182 bytes result sent to driver

22/12/19 11:52:43 INFO TaskSetManager: Starting task 96.0 in stage 0.0 (TID 96) (d0a278a3161b, executor driver, partition 96, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:43 INFO Executor: Running task 96.0 in stage 0.0 (TID 96)

22/12/19 11:52:43 INFO TaskSetManager: Finished task 85.0 in stage 0.0 (TID 85) in 5654 ms on d0a278a3161b (executor driver) (85/100)

22/12/19 11:52:43 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1862864640+19404840

22/12/19 11:52:44 INFO PythonRunner: Times: total = 7187, boot = 12, init = 10, finish = 7165

22/12/19 11:52:44 INFO Executor: Finished task 84.0 in stage 0.0 (TID 84). 2197 bytes result sent to driver

22/12/19 11:52:44 INFO TaskSetManager: Starting task 97.0 in stage 0.0 (TID 97) (d0a278a3161b, executor driver, partition 97, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:44 INFO TaskSetManager: Finished task 84.0 in stage 0.0 (TID 84) in 7198 ms on d0a278a3161b (executor driver) (86/100)

22/12/19 11:52:44 INFO Executor: Running task 97.0 in stage 0.0 (TID 97)

22/12/19 11:52:44 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1882269480+19404840

22/12/19 11:52:45 INFO PythonRunner: Times: total = 5323, boot = 5, init = 2, finish = 5316

22/12/19 11:52:45 INFO Executor: Finished task 88.0 in stage 0.0 (TID 88). 2185 bytes result sent to driver

22/12/19 11:52:45 INFO TaskSetManager: Starting task 98.0 in stage 0.0 (TID 98) (d0a278a3161b, executor driver, partition 98, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:45 INFO TaskSetManager: Finished task 88.0 in stage 0.0 (TID 88) in 5332 ms on d0a278a3161b (executor driver) (87/100)

22/12/19 11:52:45 INFO Executor: Running task 98.0 in stage 0.0 (TID 98)

22/12/19 11:52:45 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1901674320+19404840

22/12/19 11:52:46 INFO PythonRunner: Times: total = 5166, boot = 8, init = 5, finish = 5153

22/12/19 11:52:46 INFO Executor: Finished task 91.0 in stage 0.0 (TID 91). 2207 bytes result sent to driver

22/12/19 11:52:46 INFO TaskSetManager: Starting task 99.0 in stage 0.0 (TID 99) (d0a278a3161b, executor driver, partition 99, PROCESS_LOCAL, 4499 bytes) taskResourceAssignments Map()

22/12/19 11:52:46 INFO Executor: Running task 99.0 in stage 0.0 (TID 99)

22/12/19 11:52:46 INFO TaskSetManager: Finished task 91.0 in stage 0.0 (TID 91) in 5174 ms on d0a278a3161b (executor driver) (88/100)

22/12/19 11:52:46 INFO HadoopRDD: Input split: file:/opt/application/data.txt:1921079160+19404878

22/12/19 11:52:46 INFO PythonRunner: Times: total = 8057, boot = 14, init = 13, finish = 8030

22/12/19 11:52:46 INFO Executor: Finished task 86.0 in stage 0.0 (TID 86). 2205 bytes result sent to driver

22/12/19 11:52:46 INFO TaskSetManager: Finished task 86.0 in stage 0.0 (TID 86) in 8074 ms on d0a278a3161b (executor driver) (89/100)

22/12/19 11:52:46 INFO PythonRunner: Times: total = 5728, boot = -46, init = 48, finish = 5726

22/12/19 11:52:46 INFO Executor: Finished task 90.0 in stage 0.0 (TID 90). 2199 bytes result sent to driver

22/12/19 11:52:46 INFO TaskSetManager: Finished task 90.0 in stage 0.0 (TID 90) in 5736 ms on d0a278a3161b (executor driver) (90/100)

22/12/19 11:52:47 INFO PythonRunner: Times: total = 5337, boot = 1, init = 8, finish = 5328

22/12/19 11:52:47 INFO Executor: Finished task 92.0 in stage 0.0 (TID 92). 2207 bytes result sent to driver

22/12/19 11:52:47 INFO TaskSetManager: Finished task 92.0 in stage 0.0 (TID 92) in 5357 ms on d0a278a3161b (executor driver) (91/100)

22/12/19 11:52:47 INFO PythonRunner: Times: total = 8071, boot = -43, init = 56, finish = 8058

22/12/19 11:52:47 INFO Executor: Finished task 87.0 in stage 0.0 (TID 87). 2201 bytes result sent to driver

22/12/19 11:52:47 INFO TaskSetManager: Finished task 87.0 in stage 0.0 (TID 87) in 8094 ms on d0a278a3161b (executor driver) (92/100)

22/12/19 11:52:47 INFO PythonRunner: Times: total = 5015, boot = 5, init = 1, finish = 5009

22/12/19 11:52:47 INFO Executor: Finished task 94.0 in stage 0.0 (TID 94). 2205 bytes result sent to driver

22/12/19 11:52:47 INFO TaskSetManager: Finished task 94.0 in stage 0.0 (TID 94) in 5022 ms on d0a278a3161b (executor driver) (93/100)

22/12/19 11:52:47 INFO PythonRunner: Times: total = 7146, boot = 10, init = 5, finish = 7131

22/12/19 11:52:47 INFO Executor: Finished task 89.0 in stage 0.0 (TID 89). 2206 bytes result sent to driver

22/12/19 11:52:47 INFO TaskSetManager: Finished task 89.0 in stage 0.0 (TID 89) in 7154 ms on d0a278a3161b (executor driver) (94/100)

22/12/19 11:52:48 INFO PythonRunner: Times: total = 6190, boot = 9, init = 19, finish = 6162

22/12/19 11:52:48 INFO Executor: Finished task 93.0 in stage 0.0 (TID 93). 2213 bytes result sent to driver

22/12/19 11:52:48 INFO TaskSetManager: Finished task 93.0 in stage 0.0 (TID 93) in 6197 ms on d0a278a3161b (executor driver) (95/100)

22/12/19 11:52:48 INFO PythonRunner: Times: total = 5294, boot = 5, init = 5, finish = 5284

22/12/19 11:52:48 INFO Executor: Finished task 95.0 in stage 0.0 (TID 95). 2194 bytes result sent to driver

22/12/19 11:52:48 INFO TaskSetManager: Finished task 95.0 in stage 0.0 (TID 95) in 5300 ms on d0a278a3161b (executor driver) (96/100)

22/12/19 11:52:48 INFO PythonRunner: Times: total = 4877, boot = -3, init = 12, finish = 4868

22/12/19 11:52:48 INFO Executor: Finished task 96.0 in stage 0.0 (TID 96). 2200 bytes result sent to driver

22/12/19 11:52:48 INFO TaskSetManager: Finished task 96.0 in stage 0.0 (TID 96) in 4888 ms on d0a278a3161b (executor driver) (97/100)

22/12/19 11:52:49 INFO PythonRunner: Times: total = 4675, boot = 5, init = 2, finish = 4668

22/12/19 11:52:49 INFO Executor: Finished task 97.0 in stage 0.0 (TID 97). 2196 bytes result sent to driver

22/12/19 11:52:49 INFO TaskSetManager: Finished task 97.0 in stage 0.0 (TID 97) in 4684 ms on d0a278a3161b (executor driver) (98/100)

22/12/19 11:52:49 INFO PythonRunner: Times: total = 4326, boot = 3, init = 7, finish = 4316

22/12/19 11:52:49 INFO Executor: Finished task 98.0 in stage 0.0 (TID 98). 2196 bytes result sent to driver

22/12/19 11:52:49 INFO TaskSetManager: Finished task 98.0 in stage 0.0 (TID 98) in 4332 ms on d0a278a3161b (executor driver) (99/100)

22/12/19 11:52:50 INFO PythonRunner: Times: total = 3991, boot = 3, init = 3, finish = 3985

22/12/19 11:52:50 INFO TaskSetManager: Finished task 99.0 in stage 0.0 (TID 99) in 3999 ms on d0a278a3161b (executor driver) (100/100)

22/12/19 11:52:50 INFO TaskSchedulerImpl: Removed TaskSet 0.0, whose tasks have all completed, from pool

22/12/19 11:52:50 INFO DAGScheduler: ResultStage 0 (takeOrdered at /opt/application/main.py:15) finished in 64.140 s

22/12/19 11:52:50 INFO DAGScheduler: Job 0 is finished. Cancelling potential speculative or zombie tasks for this job

22/12/19 11:52:50 INFO TaskSchedulerImpl: Killing all running tasks in stage 0: Stage finished

22/12/19 11:52:50 INFO DAGScheduler: Job 0 finished: takeOrdered at /opt/application/main.py:15, took 64.222303 s

['Everybody consider article administration discussion whatever professional available.', 'Newspaper government information everybody remember prevent strategy responsibility.', 'Professional performance Democrat soldier development society successful although.', 'American sea responsibility organization itself themselves officer administration.', 'Technology environmental knowledge everybody maintain stuff expect administration.', 'Measure center democratic responsibility organization chance structure statement.', 'Identify necessary behavior individual pressure garden information international.', 'Special performance south American international receive administration describe.', 'Administration commercial challenge organization collection condition hair line.', 'Responsibility suggest environment once modern commercial international herself.']

22/12/19 11:52:50 INFO SparkContext: Invoking stop() from shutdown hook

22/12/19 11:52:50 INFO SparkUI: Stopped Spark web UI at http://d0a278a3161b:4040

22/12/19 11:52:50 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!

22/12/19 11:52:50 INFO MemoryStore: MemoryStore cleared

22/12/19 11:52:50 INFO BlockManager: BlockManager stopped

22/12/19 11:52:50 INFO BlockManagerMaster: BlockManagerMaster stopped

22/12/19 11:52:50 INFO OutputCommitCoordinator\$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!

22/12/19 11:52:50 INFO SparkContext: Successfully stopped SparkContext

22/12/19 11:52:50 INFO ShutdownHookManager: Shutdown hook called

22/12/19 11:52:50 INFO ShutdownHookManager: Deleting directory /tmp/spark-a462bf8a-fe9d-45f7-9cc5-9026a0d56505

22/12/19 11:52:50 INFO ShutdownHookManager: Deleting directory /tmp/spark-24cf905f-c636-47d3-95cc-c6427768d9de

22/12/19 11:52:50 INFO ShutdownHookManager: Deleting directory /tmp/spark-a462bf8a-fe9d-45f7-9cc5-9026a0d56505/pyspark-e329c153-b132-4053-b860-b39d29b4214f

RESULT (Топ 10 найдовших текстових коментарів):

['Everybody consider article administration discussion whatever professional available.',
'Newspaper government information everybody remember prevent strategy responsibility.',
'Professional performance Democrat soldier development society successful although.',
'American sea responsibility organization itself themselves officer administration.',
'Technology environmental knowledge everybody maintain stuff expect administration.',
'Measure center democratic responsibility organization chance structure statement.',
'Identify necessary behavior individual pressure garden information international.',
'Special performance south American international receive administration describe.',
'Administration commercial challenge organization collection condition hair line.',
'Responsibility suggest environment once modern commercial international herself.']